

Computer Graphics

Lecture 07: Ray Tracing

DR. ELVIS S. LIU


SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

SPRING 2018

A solid green horizontal bar at the bottom of the slide.

Rendering

- A process of generating an image from a 2D or 3D model, by means of computer programs




Rendering (cont.)

- “Rendering” refers to the entire process that produces color values for pixels, given a 3D representation of the scene
- Pixels correspond to rays; need to figure out the visible scene point along each ray
 - Called “hidden surface problem” in older texts
 - “Visibility” is a more modern term
 - Also, we assume (for now) a single ray per pixel
- Major algorithms: Ray casting

Rendering Methods

- Forward tracing
 - Consider every photon emitted by light source
 - Inefficient and extremely slow
 - Large number of light rays
- Backward tracing (ray tracing)
 - Trace hypothetical photo backwards from eye in particular direction
 - Slow
 - Number of light rays equals to the screen resolution
- Surface rendering
 - What light shines on surface?
 - How does material interact with light?
 - What part of result is visible to eye?



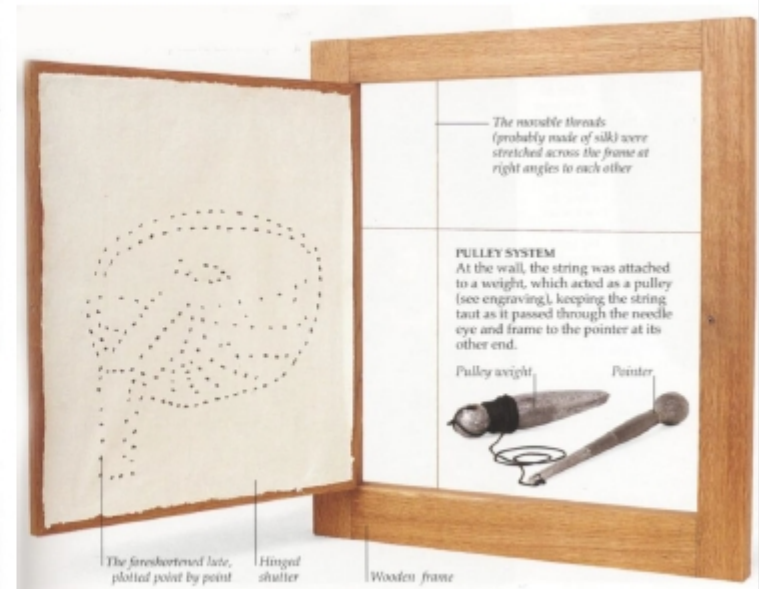

Origins of Ray Tracing

LECTURE 07: RAY TRACING

A solid green horizontal bar at the bottom of the slide.

Dürer's Ray Casting Machine

- Generalising from Albrecht Dürer's wood cut (16th Century) showing perspective projection
- Durer: Record string intersection from center of projection (eye), to nearest object as points on a 2D plane




Dürer's Ray Casting Machine (cont.)

- Points created are perspective projection of 3D object onto 2D plane – our pixels
- Can think of first starting with sample points on objects and then drawing ray OR starting with ray thru pixel center (or sample points within supersampled pixel)



What is a Raytracer?

- A finite back-mapping of rays from camera (eye) through each sample (pixel or subpixel) to objects in scene
 - Avoids forward solution of having to sample from an infinite number of rays from light sources, not knowing which will be important for PoV
- Each pixel represents either:
 - A ray intersection with an object/light in scene
 - No intersection
- A ray traced scene is a “virtual photo” comprised of many samples on film plane
- Generalising from one ray, millions of rays are shot from eye, one through each point on film plane



Ray Tracing Fundamentals

- Generate primary ray
 - shoot rays from eye through sample points on film plane
 - sample point is typically center of a pixel, but alternatively supersample pixels (recall supersampling from Image Processing IV)
- Ray-object intersection
 - find first object in scene that ray intersects with (if any)
 - solves VSD/HSR problem – use parametric line equation for ray, so smallest t value


Ray Tracing Fundamentals (cont.)

- Calculate lighting (i.e., colour)
 - Use illumination model to determine direct contribution from light sources (light rays)
 - Reflective objects recursively generate secondary rays (indirect contribution) that also contribute to color; RRT (Recursive Ray Tracing) only uses specular reflection rays because they are easy to compute and specular reflections are typically brighter than average
 - Sum of contributions determines color of sample point
 - No diffuse reflection rays → RRT is limited approximation to global illumination
- Finesse need for shading rule – ray-trace for lighting equation evaluation at each sample point

Ray Casting vs. Ray Tracing

- Ray Casting: eye rays only
- Ray Tracing: consider also secondary rays which are used for testing shadows, doing reflections, refractions, etc.

Ray Tracing Examples



Ray Tracing Examples (cont.)




Ray Generation

LECTURE 07: RAY TRACING


Simple Ray Casting

- For every pixel
 - Construct a ray from the eye
 - For every object in the scene
 - Find intersection with the ray
 - Keep if closest




Ray Origin

- Consider the geometry of the problem in normalised world space with canonical perspective frustum (i.e., do not apply perspective transformation)
- Start a ray from “eye point” \mathbf{P}
- Shoot ray in some direction \mathbf{d} from \mathbf{P} toward a point in film plane (a rectangle in the \mathbf{u} - \mathbf{v} plane in the camera’s \mathbf{uvw} space) whose color we want to know
- Points along ray have form $\mathbf{P} + t\mathbf{d}$ where
 - \mathbf{P} is ray’s base point: camera’s eye
 - \mathbf{d} is unit vector direction of ray
 - t is a non-negative real number
- “Eye point” \mathbf{P} is center of projection in perspective view volume (view frustum)




Ray Direction

- Start with 2D screen-space sample point ((sub)pixel)
- To create a ray from eye point through film plane, 2D screen-space point must be converted into a 3D point on film plane. Then we can get the direction vector by subtracting the COP from the 3D screen space point
- Note that ray generated will be intersecting objects in normalised world space coordinate system before the perspective transformation




canonical frustum in normalized world coordinates
Any plane $z = k$, $-1 \leq k < 0$ can be the film plane




Normalised world space

Ray Direction (cont.)

- Any plane orthogonal to look vector is a convenient film plane (plane $z = k$ in canonical frustum)
- Choose a film plane and then create a function that maps screen integer (u, v) space points onto it to floats
 - what's a good plane to use? Try the far clipping plane ($z = -1$)
 - to convert, scale integer screen-space coordinates into floats between -1 and 1 for x and y , $z = -1$




canonical frustum in normalized world coordinates
Any plane $z = k$, $-1 \leq k < 0$ can be the film plane



Normalised world space


Ray Direction (cont.)

- Once have a 3D point on the film plane, need to transform to pre-normalisation world space where objects are
 - Why? illumination model prefers intersection point to be in world-space (less work than normalising lights)
 - Make a direction vector from eye point P (at center of projection) to 3D point on film plane
 - Need this vector to be in world-space in order to intersect with original object in pre-normalisation world space




Ray Direction (cont.)

- Normalising transformation maps world-space points to points in the canonical view volume
 - translate to the origin; orient so Look points down $-Z$, Up along Y ; scale x and y to adjust viewing angles to 45° , scale z : $[-1, 0]$; x, y : $[-1, 1]$
- How do we transform a point from the canonical view volume back to untransformed world space?
 - apply the inverse of the normalising transformation: Viewing Transformation
 - Note: not same as viewing transform you use in OpenGL (e.g., ModelView matrix)



Summary

- Start the ray at center of projection (“eye point”)
- Map 2D integer screen-space point onto 3D film plane in normalised frustum
 - scale x, y to fit between -1 and 1
 - set z to -1 so points lie on the far clip plane
- Transform 3D film plane point (mapped (sub)pixel) to an untransformed world-space point
 - need to undo normalising transformation (i.e., viewing transformation)
- Construct the direction vector
 - a point minus a point is a vector
 - $\text{direction} = (\text{world-space point (mapped pixel)}) - (\text{eye point (in untransformed world space)})$



Ray-Object Intersection

LECTURE 07: RAY TRACING


A solid green horizontal bar spanning the width of the slide at the bottom.

Implicit Objects

- If an object is defined implicitly by a function f , where $f(Q) = 0$ IFF Q is a point on the surface of the object, then ray-object intersections are relatively easy to compute
 - Many objects can be defined implicitly
 - Implicit functions provide potentially infinite resolution
 - Tessellating these objects is more difficult than using the implicit functions directly
- For example, a circle of radius R is an implicit object in a plane, with equation:
 - $f(x,y) = x^2 + y^2 - R^2$
 - point (x,y) is on the circle when $f(x,y) = 0$
- A sphere of radius R in 3-space:
 - $f(x,y,z) = x^2 + y^2 + z^2 - R^2$

Implicit Objects (cont.)

- At what points (if any) does the ray intersect an object?
- Points on a ray have form $\mathbf{P} + t\mathbf{d}$ where t is any non-negative real number
- A surface point Q lying on an object has the property that $f(Q) = 0$
- Combining, we want to know “For which values of t is $f(\mathbf{P} + t\mathbf{d}) = 0$?”
- We are solving a system of simultaneous equations in x, y (in 2D) or x, y, z (in 3D)



Explicit vs. Implicit

- Ray equation is explicit $P(t) = R_o + t \times R_d$
 - Parametric
 - Generates points
 - Hard to verify that a point is on the ray
- Plane equation is implicit $H(P) = n \cdot P + D = 0$
 - Solution of an equation
 - Does not generate points
 - Verifies that a point is on the plane
- What about implicit plane and explicit ray?

An Explicit Example – 2D Ray-circle Intersection


- Consider the eye-point $\mathbf{P} = (-3, 1)$, the direction vector $\mathbf{d} = (0.8, -0.6)$ and the unit circle: $f(x,y) = x^2 + y^2 - R^2$
- A typical point of the ray is: $\mathbf{Q} = \mathbf{P} + t\mathbf{d} = (-3,1) + t(0.8, -0.6) = (-3+0.8t, 1-0.6t)$
- Plugging this into the equation of the circle:
 $f(\mathbf{Q}) = f(-3+0.8t, 1-0.6t) = (-3+0.8t)^2 + (1-0.6t)^2 - 1$
- Expanding, we get: $9 - 4.8t + 0.64t^2 + 1 - 1.2t + .36t^2 - 1$
- Setting this equal to zero, we get: $t^2 - 6t + 9 = 0$

An Explicit Example – 2D Ray-circle Intersection (cont.)

- Using the quadratic formula: $roots = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$
- We get: $t = \frac{6 \pm \sqrt{36 - 36}}{2}, \quad t = 3, 3$
- Because we have a root of multiplicity 2, ray intersects circle at only one point (i.e., it's tangent to the circle)
- Use discriminant $D = b^2 - 4ac$ to quickly determine if true intersection:
 - if $D < 0$, imaginary roots; no intersection
 - if $D = 0$, double root; ray is tangent
 - if $D > 0$, two real roots; ray intersects circle at two points
- Smallest non-negative real t represents the intersection nearest to eye-point

An Explicit Example – 2D Ray-circle Intersection (cont.)

- Generalising:
 - We can take an arbitrary implicit surface with equation $f(Q) = 0$, A ray $\mathbf{P} + t\mathbf{d}$, and plug the latter into the former:
 - $f(\mathbf{P} + t\mathbf{d}) = 0$
- The result, after some algebra, is an equation with t as unknown
- We then solve for t , analytically or numerically



Cylindrical Objects – Multiple Conditions

- For cylindrical objects, the implicit equation
 - $f(x, y, z) = x^2 + z^2 - 1 = 0$
 - In 3D-space defines an infinite cylinder of unit radius running along the y-axis
- Usually, it's more useful to work with finite objects
 - e.g. a unit cylinder truncated with the limits:
 - Cylinder body: $x^2 + z^2 - 1 = 0, -1 \leq y \leq 1$
- But how do we define cylinder “caps” as implicit equations?
- The caps are the insides of the cylinder at the cylinder's y extrema (or rather, a circle)
 - Cylinder caps
 - top: $x^2 + z^2 - 1 \leq 0, y = 1$
 - bottom: $x^2 + z^2 - 1 \leq 0, y = -1$

Cylinder Pseudocode

- Solve in a case-by-case approach

Ray_inter_finite_cylinder(P,d):

t1,t2 = ray_inter_infinite_cylinder(P,d)

// Check for intersection with infinite cylinder

compute $P + t1*d$, $P + t2*d$

if $y > 1$ or $y < -1$ for t1 or t2: toss it

// If intersection, is it between cylinder caps?

t3 = ray_inter_plane(plane $y = 1$)

// Check for an intersection with the top cap

Compute $P + t3*d$

if $x^2 + z^2 > 1$: toss out t3

// If it intersects, is it within cap circle?

t4 = ray_inter_plane(plane $y = -1$)

// Check intersection with bottom cap

Compute $P + t4*d$

if $x^2 + z^2 > 1$: toss out t4



// If it intersects, is it within cap circle?

Of all the remaining t's (t1 – t4), select the smallest non-negative one.

If none remain, ray does not intersect cylinder


Implicit Surface Strategy Summary

- Substitute ray ($P + td$) into implicit surface equations and solve for t
 - smallest non-negative t -value is from the closest surface you see from eye point
- For complicated objects (not defined by a single equation), write out a set of equalities and inequalities and then code individual surfaces as cases
- Latter approach can be generalized cleverly to handle all sorts of complex combinations of objects
 - constructive Solid Geometry (CSG), where objects are stored as a hierarchy of primitives and 3-D set operations (union, intersection, difference) – don't have to evaluate the hierarchy to raytrace!
 - “blobby objects”, which are implicit surfaces defined by sums of implicit equations ($F(x,y,z)=0$)




Ray-Triangle Intersection

- Motivation: 3D objects are usually represented by triangle meshes
- Use ray-plane intersection followed by in-triangle test
- Or try to be smarter
 - Use barycentric coordinates




Barycentric Definition of a Plane

- A (non-degenerate) triangle (a,b,c) defines a plane
- Any point P on this plane can be written as
 - $P(\alpha,\beta,\gamma) = \alpha \mathbf{a} + \beta \mathbf{b} + \gamma \mathbf{c}$,
 - with $\alpha + \beta + \gamma = 1$




Barycentric Coordinates

- Since $\alpha + \beta + \gamma = 1$, we can write
 - $\alpha = 1 - \beta - \gamma$
- $\mathbf{P}(\alpha, \beta, \gamma) = \alpha \mathbf{a} + \beta \mathbf{b} + \gamma \mathbf{c}$
- $\mathbf{P}(\beta, \gamma) = (1 - \beta - \gamma) \mathbf{a} + \beta \mathbf{b} + \gamma \mathbf{c}$
 - $= \mathbf{a} + \beta(\mathbf{b} - \mathbf{a}) + \gamma(\mathbf{c} - \mathbf{a})$




Barycentric Definition of a Plane

- $\mathbf{P}(\alpha, \beta, \gamma) = \alpha \mathbf{a} + \beta \mathbf{b} + \gamma \mathbf{c}$
 - With $\alpha + \beta + \gamma = 1$
- Is it explicit or implicit?
- \mathbf{P} is the **barycenter**, the single point upon which the triangle would balance if weights of size α , β , & γ are placed on points \mathbf{a} , \mathbf{b} & \mathbf{c}




Barycentric Definition of a Plane (cont.)

- $\mathbf{P}(\alpha, \beta, \gamma) = \alpha \mathbf{a} + \beta \mathbf{b} + \gamma \mathbf{c}$
 - With $\alpha + \beta + \gamma = 1$ parameterises the entire plane
- If we require in addition that $\alpha, \beta, \gamma \geq 0$, we get just the triangle
 - Note that with $\alpha + \beta + \gamma = 1$ this implies $0 \leq \alpha \leq 1$ & $0 \leq \beta \leq 1$ & $0 \leq \gamma \leq 1$
 - Verify:
 - $\alpha = 0 \Rightarrow \mathbf{P}$ lies on line $\mathbf{b}-\mathbf{c}$
 - $\alpha, \beta = 0 \Rightarrow \mathbf{P} = \mathbf{c}$
 - etc.



Barycentric Definition of a Plane (cont.)

- $\mathbf{P}(\alpha, \beta, \gamma) = \alpha \mathbf{a} + \beta \mathbf{b} + \gamma \mathbf{c}$
 - With $\alpha + \beta + \gamma = 1$ parameterises the entire plane
- Condition to be barycentric coordinates:
 - $\alpha + \beta + \gamma = 1$
- Condition to be inside the triangle
 - $\alpha, \beta, \gamma \geq 0$




How to Compute α , β , γ ?

- Ratio of opposite sub-triangle area to total area


- $\alpha = A_a/A$
 - $\beta = A_b/A$
 - $\gamma = A_c/A$

- Use signed areas for points outside the triangle



Intersection with Barycentric Triangle

- Again, set ray equation equal to barycentric equation
- $\mathbf{P}(t) = \mathbf{P}(\beta, \gamma)$
- $\mathbf{R}_o + t \mathbf{R}_d = \mathbf{a} + \beta(\mathbf{b}-\mathbf{a}) + \gamma(\mathbf{c}-\mathbf{a})$
- Intersection if
 - $\beta + \gamma \leq 1$, and
 - $\beta \geq 0$, and
 - $\gamma \geq 0$



Intersection with Barycentric Triangle (cont.)

- $\mathbf{R}_o + t * \mathbf{R}_d = \mathbf{a} + \beta(\mathbf{b}-\mathbf{a}) + \gamma(\mathbf{c}-\mathbf{a})$

- $R_{ox} + tR_{dx} = a_x + \beta(b_x - a_x) + \gamma(c_x - a_x)$

- $R_{oy} + tR_{dy} = a_y + \beta(b_y - a_y) + \gamma(c_y - a_y)$

- $R_{oz} + tR_{dz} = a_z + \beta(b_z - a_z) + \gamma(c_z - a_z)$

- Regroup and write in matrix form $Ax=b$


- $$\begin{bmatrix} a_x - b_x & a_x - c_x & R_{dx} \\ a_y - b_y & a_y - c_y & R_{dy} \\ a_z - b_z & a_z - c_z & R_{dz} \end{bmatrix} \begin{bmatrix} \beta \\ \gamma \\ t \end{bmatrix} = \begin{bmatrix} a_x - R_{ox} \\ a_y - R_{oy} \\ a_z - R_{oz} \end{bmatrix}$$

Barycentric Intersection Advantages

- Efficient
- Stores no plane equation
- Get the barycentric coordinates for free
- Useful for interpolation, texture mapping

Barycentric Interpolation

- Values v_1, v_2, v_3 defined at a, b, c
 - Colors, normal, texture coordinates, etc.
- $\mathbf{P}(\alpha, \beta, \gamma) = \alpha \mathbf{a} + \beta \mathbf{b} + \gamma \mathbf{c}$ is the point...
- $\mathbf{v}(\alpha, \beta, \gamma) = \alpha v_1 + \beta v_2 + \gamma v_3$ is the barycentric interpolation of v_1, v_2, v_3 at point \mathbf{P}
 - Sanity check: $\mathbf{v}(1,0,0) = v_1$, etc.
- i.e, once you know α, β, γ you can interpolate values using the same weights.
 - Convenient!



Summary of Ray Casting (put it all together)

$P = \text{eyePt}$

for each sample of image:

 Compute d

for each object:

 Take world space intersection point $P+td$

 Convert to object space, $+t$

 Use implicit formula to calculate t for object

 Store object space intersection for later normal calculation

 Select object with smallest non-negative t -value (visible object)

 Compute object space normal at object space intersection

 Transform object space normal to world space

 Use normal and intersection point (both in world space) for lighting computation

Shadows

LECTURE 07: RAY TRACING

Equation of Shadows


- Each light in the scene contributes to the color and intensity of a surface element...

$$objectIntensity_{\lambda} = ambient + \sum_{light=1}^{numLights} attenuation \cdot lightIntensity_{\lambda} \cdot [diffuse + specular]$$

- If and only if light source reaches the object
 - could be occluded/obstructed by other objects in scene
 - could be self-occluding

Algorithm of Shadows


- Construct a ray from the surface intersection to each light
- Check if light is first object intersected
 - if first object intersected is the light, count light's full contribution
 - if first object intersected is not the light, do not count (ignore) light's contribution
 - this method generates hard shadows; soft shadows are harder to compute (must sample)
- What about transparent or specular (reflective) objects? Such lighting contributions are the beginning of global illumination \Rightarrow need recursive ray tracing



Recursive Ray Tracing


LECTURE 07: RAY TRACING

Recursive Ray Tracing



Simulating global lighting effects

- By recursively casting new rays into the scene, we can look for more information
- Start from point of intersection with object
- We'd like to send rays in all directions, but that's too hard/computationally taxing
- Instead, just send rays in directions likely to contribute most:
 - toward lights (blockers to lights create shadows for those lights)
 - specular bounce off other objects to capture specular inter-object reflections
 - use ambient hack to capture diffuse inter-object reflection
 - refractive rays through object (transparency)



Trace Secondary Rays at Intersections

- Light: trace a ray to each light source. If light source is blocked by an opaque object, it does not contribute to lighting
- Specular reflection: trace reflection ray (i.e., about normal vector at surface intersection)
- Refractive transmission/transparency: trace refraction ray (following Snell's law)
- Recursively spawn new light, reflection, and refraction rays at each intersection until contribution negligible or some max recursion depth is reached
- Limitations
 - recursive inter-object reflection is strictly specular
 - diffuse inter-object reflection is handled by other techniques

Lighting Equation for Recursive Ray Tracing

- New lighting equation (Phong lighting + specular reflection + transmission):

$$I_{\lambda} = \underbrace{L_{a\lambda} k_a O_{a\lambda}}_{\text{ambient}} + \sum_{\text{lights}} f_{\text{att}} L_{p\lambda} \left[\underbrace{k_d O_{d\lambda} (\vec{n} \bullet \vec{l})}_{\text{diffuse}} + \underbrace{k_s O_{s\lambda} (\vec{r} \bullet \vec{v})^n}_{\text{specular}} \right] + \underbrace{k_s O_{s\lambda} I_{r\lambda}}_{\text{reflected}} + \underbrace{k_t O_{t\lambda} I_{t\lambda}}_{\text{transmitted}}$$

recursive rays

- I is the total color at a given point (lighting + specular reflection + transmission, λ subscript for each r,g,b)
- Its presence in the transmitted and reflected terms implies recursion
- L is the light intensity; L_p is the intensity of a point light source
- k is the attenuation coefficient for the object material (ambient, diffuse, specular, etc.)
- O is the object color
- f_{att} is the attenuation function for distance
- n is the normal vector at the object surface

Lighting Equation for Recursive Ray Tracing (cont.)

- New lighting equation (Phong lighting + specular reflection + transmission):

$$I_{\lambda} = \underbrace{L_{a\lambda} k_a O_{a\lambda}}_{\text{ambient}} + \sum_{\text{lights}} f_{\text{att}} L_{p\lambda} \left[\underbrace{k_d O_{d\lambda} (\vec{n} \bullet \vec{l})}_{\text{diffuse}} + \underbrace{k_s O_{s\lambda} (\vec{r} \bullet \vec{v})^n}_{\text{specular}} \right] + \underbrace{k_s O_{s\lambda} I_{r\lambda}}_{\text{reflected}} + \underbrace{k_t O_{t\lambda} I_{t\lambda}}_{\text{transmitted}}$$

recursive rays

- l is the vector to the light
 - r is the reflected light vector
 - v is the vector from the eye point (view vector)
 - n is the specular exponent
 - note: intensity from recursive rays calculated with the same lighting equation at the intersection point
 - light sources contribute specular and diffuse lighting
- Note: single rays of light do not attenuate with distance; purpose of f_{att} is to simulate diminishing intensity per unit area as function of distance for point lights (typically an inverse quadratic polynomial)

Some Old Videos

- <https://www.youtube.com/watch?v=sg8YwRNA5j8>

Some Old Videos (cont.)

- https://www.youtube.com/watch?v=b_UqzLBFz4Y

Transparent Surfaces

LECTURE 07: RAY TRACING

Non-Refractive Transparency


- For a partially transparent polygon

$$I_{\lambda} = (1 - k_{t1})I_{\lambda1} + k_{t1}I_{\lambda2}$$

k_{t1} = transmittance of polygon 1
(0 = opaque; 1 = transparent)

$I_{\lambda1}$ = intensity calculated for polygon 1

$I_{\lambda2}$ = intensity calculated for polygon 2




Refractive Transparency

- Model the way light bends at interfaces with Snell's Law

$$\sin \theta_r = \frac{\sin \theta_i \eta_{i\lambda}}{\eta_{r\lambda}}$$

$\eta_{i\lambda}$ = index of refraction of medium 1

$\eta_{r\lambda}$ = index of refraction of medium 2




Sampling

LECTURE 07: RAY TRACING


Choosing Samples

- This example samples once per pixel. This generates images similar to this:
- We have a clear case of the jaggies
- Can we do better?




Choosing Samples (cont.)

- In the simplest case, sample points are chosen at pixel centers
- For better results, supersamples can be chosen (called supersampling)
 - e.g., at corners of pixel as well as at center
- Even better techniques do adaptive sampling: increase sample density in areas of rapid change (in geometry or lighting)
- With stochastic sampling, samples are taken probabilistically (recall Image Processing IV slides)
 - Actually converges on “correct” answer faster than regularly spaced sampling
- For fast results, we can subsample: fewer samples than pixels
 - take as many samples as time permits
 - beam tracing: track a bundle of neighboring rays together
- How do we convert samples to pixels? Filter to get weighted average of all the samples per pixel!
 - Instead of sampling one point, sample within a region to create a better approximation



vs.



Supersampling Example

WITH SUPERSAMPLING



WITHOUT SUPERSAMPLING

